

Alphanumeric recognition

In [1]:

```
import pickle
import tensorflow as tf
from tensorflow.python import keras
from tensorflow.python.keras.datasets import mnist
from tensorflow.python.keras.models import Sequential
from tensorflow.python.keras.layers import Dense, Dropout, Input, Flatten, Activation
from tensorflow.python.keras.layers import Conv2D, MaxPool2D
from tensorflow.python.keras.optimizers import RMSprop
import matplotlib.pyplot as plt
import numpy as np
import h5py
import random
import sys

from sklearn.model_selection import train_test_split
```

Unpickling data

In [2]:

```
with open("../train.pkl", "rb") as infile:
    (X, y) = pickle.load(infile)
```

In [3]:

```
X = X.reshape(30134, 56, 56, 1)
y = y.ravel()
```

Look at the data

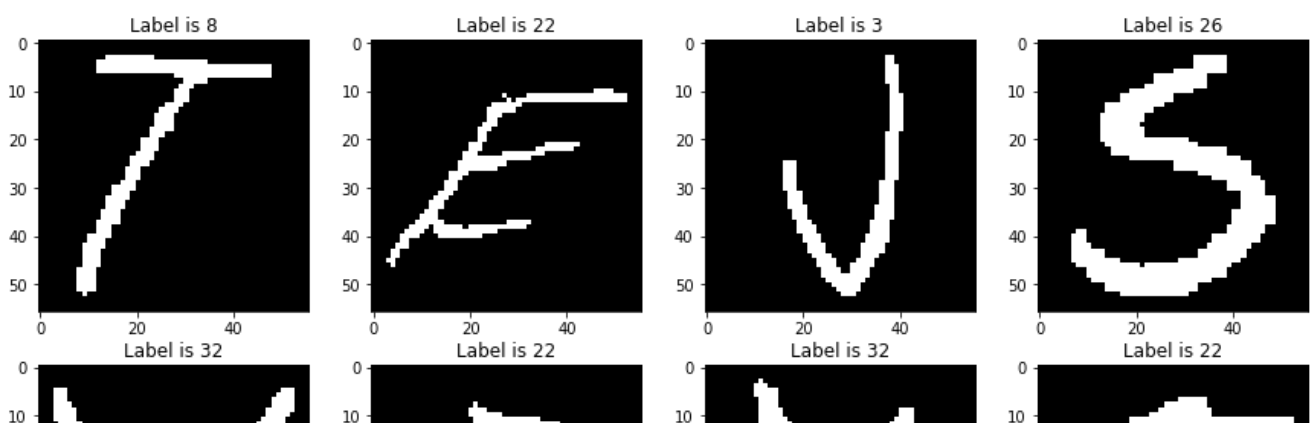
Data contains 56x56 written letters and number. Below random sample from data is shown.


In [4]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

random_indices = np.random.randint(0, X_train.shape[0], size=16)

fig = plt.figure(figsize=(15, 15))
for i, ridx in enumerate(random_indices):
    fig.add_subplot(4, 4, i+1)
    plt.title('Label is {label}'.format(label=y_train[ridx]))
    plt.imshow(X_train[ridx, :, :, 0], cmap='gray')
plt.show()
```





In [5]:

```
num_classes = 36

X_train = X_train.astype('uint8')
X_test = X_test.astype('uint8')

X_train_fc = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test_fc = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

y_train_1h = keras.utils.to_categorical(y_train, num_classes)
y_test_1h = keras.utils.to_categorical(y_test, num_classes)
```

Fully connected model

5x Fully Connected Layer

ReLU activation is used (except from output layer where softmax is used).

In [6]:

```
epochs = 30

model_fc = Sequential()
model_fc.add(Dropout(0.2, input_shape=(3136,)))
model_fc.add(Dense(700, activation='relu'))
model_fc.add(Dropout(0.5))
model_fc.add(Dense(700, activation='relu'))
model_fc.add(Dropout(0.5))
model_fc.add(Dense(500, activation='relu'))
model_fc.add(Dropout(0.5))
model_fc.add(Dense(500, activation='relu'))
model_fc.add(Dense(num_classes, activation='softmax'))

model_fc.compile(loss='categorical_crossentropy',
                  optimizer=RMSprop(),
                  metrics=['accuracy'])
```

In [7]:

```
print(model_fc.summary())
```

Layer (type)	Output Shape	Param #
dropout_1 (Dropout)	(None, 3136)	0
dense_1 (Dense)	(None, 700)	2195900
dropout_2 (Dropout)	(None, 700)	0
dense_2 (Dense)	(None, 700)	490700
dropout_3 (Dropout)	(None, 700)	0
dense_3 (Dense)	(None, 500)	350500
dropout_4 (Dropout)	(None, 500)	0
dense_4 (Dense)	(None, 500)	250500
dense_5 (Dense)	(None, 36)	18036

=====
Total params: 3,305,636
Trainable params: 3,305,636
Non-trainable params: 0
=====
None

In [10]:

```
EStop = keras.callbacks.EarlyStopping(monitor='val_loss', patience=2)
```

```
model_fc.fit(X_train_fc, y_train_1h,  
            batch_size=128,  
            epochs=epochs,  
            verbose=1,  
            validation_data=(X_test_fc, y_test_1h),  
            callbacks=[EStop])  
score_fc = model_fc.evaluate(X_test_fc, y_test_1h, verbose=0)
```

```
print('Test loss: {:.6f}'.format(score_fc[0]))  
print('Test accuracy: {:.2f} %'.format(score_fc[1]*100))
```

Train on 24107 samples, validate on 6027 samples

Epoch 1/30

24107/24107 [=====] - 13s 546us/step - loss: 0.5486 - acc: 0.8342 - val_loss: 0.4278 - val_acc: 0.8578

Epoch 2/30

24107/24107 [=====] - 11s 451us/step - loss: 0.4887 - acc: 0.8535 - val_loss: 0.4008 - val_acc: 0.8752

Epoch 3/30

24107/24107 [=====] - 11s 454us/step - loss: 0.4592 - acc: 0.8632 - val_loss: 0.4028 - val_acc: 0.8580

Epoch 4/30

24107/24107 [=====] - 12s 481us/step - loss: 0.4363 - acc: 0.8719 - val_loss: 0.3983 - val_acc: 0.8737

Epoch 5/30

24107/24107 [=====] - 11s 468us/step - loss: 0.4206 - acc: 0.8787 - val_loss: 0.3705 - val_acc: 0.8761

Epoch 6/30

24107/24107 [=====] - 12s 510us/step - loss: 0.4126 - acc: 0.8834 - val_loss: 0.3580 - val_acc: 0.8917

Epoch 7/30

24107/24107 [=====] - 13s 519us/step - loss: 0.3828 - acc: 0.8890 - val_loss: 0.3612 - val_acc: 0.8940

Epoch 8/30

24107/24107 [=====] - 11s 472us/step - loss: 0.3657 - acc: 0.8952 - val_loss: 0.3750 - val_acc: 0.8766

Test loss: 0.375020

Test accuracy: 87.66 %

Convolution model

Model contains: 1xConvolution Layer, 3x(Convolution +Max Pooling), 2x Fully Connected.

ReLU activation is used (except from output layer where softmax is used).

In [11]:

```
epochs = 30

model_cnn = Sequential()
model_cnn.add(Dropout(0.2, input_shape=(56,56,1)))
model_cnn.add(Conv2D(filters=16,kernel_size=(3,3),data_format="channels_last",input_shape=(56,56, 1,)))
model_cnn.add(Activation("relu"))
model_cnn.add(Conv2D(filters=32,kernel_size=(3,3),data_format="channels_last"))
model_cnn.add(Activation("relu"))
model_cnn.add(MaxPool2D(pool_size=(2,2)))
model_cnn.add(Dropout(0.5))
model_cnn.add(Conv2D(filters=64,kernel_size=(3,3),data_format="channels_last"))
model_cnn.add(Activation("relu"))
model_cnn.add(MaxPool2D(pool_size=(2,2),data_format="channels_last"))
model_cnn.add(Dropout(0.5))
model_cnn.add(Conv2D(filters=128,kernel_size=(3,3),data_format="channels_last"))
model_cnn.add(Activation("relu"))
model_cnn.add(MaxPool2D(pool_size=(3,3),data_format="channels_last"))
model_cnn.add(Flatten())
model_cnn.add(Dense(500, activation='relu'))
model_cnn.add(Dropout(0.5))
model_cnn.add(Dense(num_classes, activation='softmax'))

model_cnn.compile(loss='categorical_crossentropy',
                  optimizer=RMSprop(),
                  metrics=['accuracy'])
```

In [12]:

```
print(model_cnn.summary())
```

Layer (type)	Output Shape	Param #
=====		
dropout_5 (Dropout)	(None, 56, 56, 1)	0
conv2d_1 (Conv2D)	(None, 54, 54, 16)	160
activation_1 (Activation)	(None, 54, 54, 16)	0
conv2d_2 (Conv2D)	(None, 52, 52, 32)	4640
activation_2 (Activation)	(None, 52, 52, 32)	0
max_pooling2d_1 (MaxPooling2	(None, 26, 26, 32)	0
dropout_6 (Dropout)	(None, 26, 26, 32)	0
conv2d_3 (Conv2D)	(None, 24, 24, 64)	18496
activation_3 (Activation)	(None, 24, 24, 64)	0
max_pooling2d_2 (MaxPooling2	(None, 12, 12, 64)	0
dropout_7 (Dropout)	(None, 12, 12, 64)	0
conv2d_4 (Conv2D)	(None, 10, 10, 128)	73856
activation_4 (Activation)	(None, 10, 10, 128)	0
max_pooling2d_3 (MaxPooling2	(None, 3, 3, 128)	0
flatten_1 (Flatten)	(None, 1152)	0

dense_6 (Dense)	(None, 500)	576500
dropout_8 (Dropout)	(None, 500)	0
dense_7 (Dense)	(None, 36)	18036
=====		
Total params: 691,688		
Trainable params: 691,688		
Non-trainable params: 0		
None		

In [14]:

```
model_cnn.fit(X_train, y_train_1h,
              batch_size=128,
              epochs=epochs,
              verbose=1,
              validation_data=(X_test, y_test_1h),
              callbacks=[EStop])
score_cnn = model_cnn.evaluate(X_test, y_test_1h, verbose=0)

print('Test loss: {:.6f}'.format(score_cnn[0]))
print('Test accuracy: {:.2f} %'.format(score_cnn[1]*100))
```

```
Train on 24107 samples, validate on 6027 samples
Epoch 1/30
24107/24107 [=====] - 141s 6ms/step - loss: 1.1392 - acc: 0.6809 - val_loss: 0.6200 - val_acc: 0.8485
Epoch 2/30
24107/24107 [=====] - 142s 6ms/step - loss: 0.4830 - acc: 0.8479 - val_loss: 0.4406 - val_acc: 0.8966
Epoch 3/30
24107/24107 [=====] - 140s 6ms/step - loss: 0.3703 - acc: 0.8841 - val_loss: 0.5183 - val_acc: 0.9069
Epoch 4/30
24107/24107 [=====] - 140s 6ms/step - loss: 0.3139 - acc: 0.8992 - val_loss: 0.3717 - val_acc: 0.9209
Epoch 5/30
24107/24107 [=====] - 145s 6ms/step - loss: 0.2849 - acc: 0.9086 - val_loss: 0.2850 - val_acc: 0.9204
Epoch 6/30
24107/24107 [=====] - 145s 6ms/step - loss: 0.2558 - acc: 0.9175 - val_loss: 0.2669 - val_acc: 0.9215
Epoch 7/30
24107/24107 [=====] - 150s 6ms/step - loss: 0.2376 - acc: 0.9213 - val_loss: 0.2824 - val_acc: 0.9273
Epoch 8/30
24107/24107 [=====] - 162s 7ms/step - loss: 0.2233 - acc: 0.9257 - val_loss: 0.2837 - val_acc: 0.9280
Test loss: 0.283652
Test accuracy: 92.80 %
```

Choosing a model

Convolutional model offers 41.7 % error reduction on development set. It also has fewer parameters.

Misclassification error reduction is:

$$\text{reduction} = \frac{\text{error}_0 - \text{error}_1}{\text{error}_0}$$

In [15]:

```
print('Error reduction: {:.2f} %'.format((score_cnn[1]-score_fc[1])/(1-score_fc[1])*100))

model_cnn.save("prediction_model.h5")
```

Error reduction: 41.67 %

Short summary

Two simple neural networks architectures were proposed for the problem:

- a) Fully connected - data is treated as vector
- b) Convolutional - data is treated as images with height and width

For both architectures I have used dropout (0.2 in input, 0.5 in hidden layers) to regularize data during training and early stopping to find optimal time to stop the training. During training validation loss was tracked. Those measures prevent overfitting.

I have used RMSprop as gradient optimization schedule. Activation in most layers is ReLU.

Differences in architectures

In convolutional architecture I use 2D Convolution as input layer 3 sets of 2D Convolution -> ReLU -> Max Pooling layers. Convolution layers allow to better "understand" spatial data. It achieves 41.7 % error reduction on validation set. Moreover convolutional architecture has around 5 times less parameters.

Technologies

Keras model with tensorflow backend was used for the problem.